

# Analyzing 2022 NFL 4th Down Success by Team

Jack Miller and Alex Jackson

## Introduction

Our project examines NFL 4th down decision making using play-by-play data dating back to 1999. Our goal is to build three models that predict expected win probability added for each of the three decisions coaches have on fourth downs: kick a field goal, punt it, or go for it.

## Exploratory Data Analysis

Our data comes from the `nflfastR` package which contains data on every play dating back to 1999.

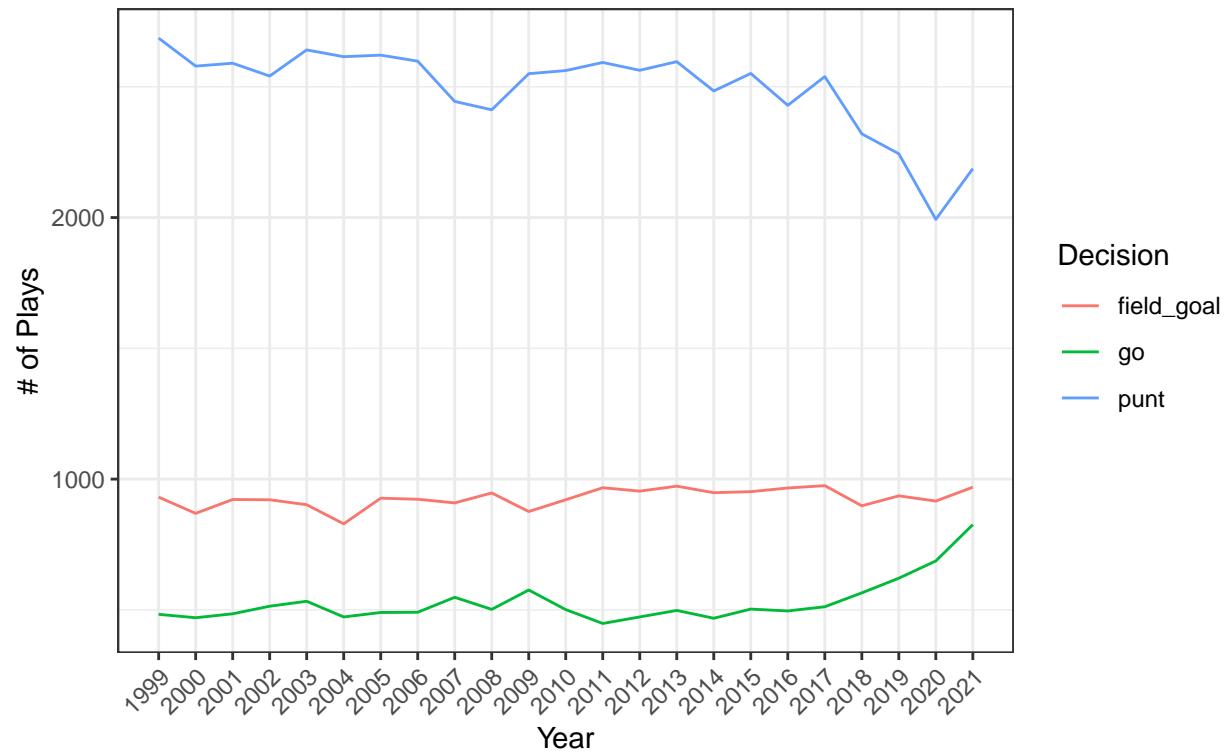
```
data_main %>%
  group_by(play_type) %>%
  summarize(count = n(), perc = n()*100/nrow(data_main)) %>%
  kable(digits = 1)
```

play_type	count	perc
field_goal	22037	23.5
go	12672	13.5
punt	58923	62.9

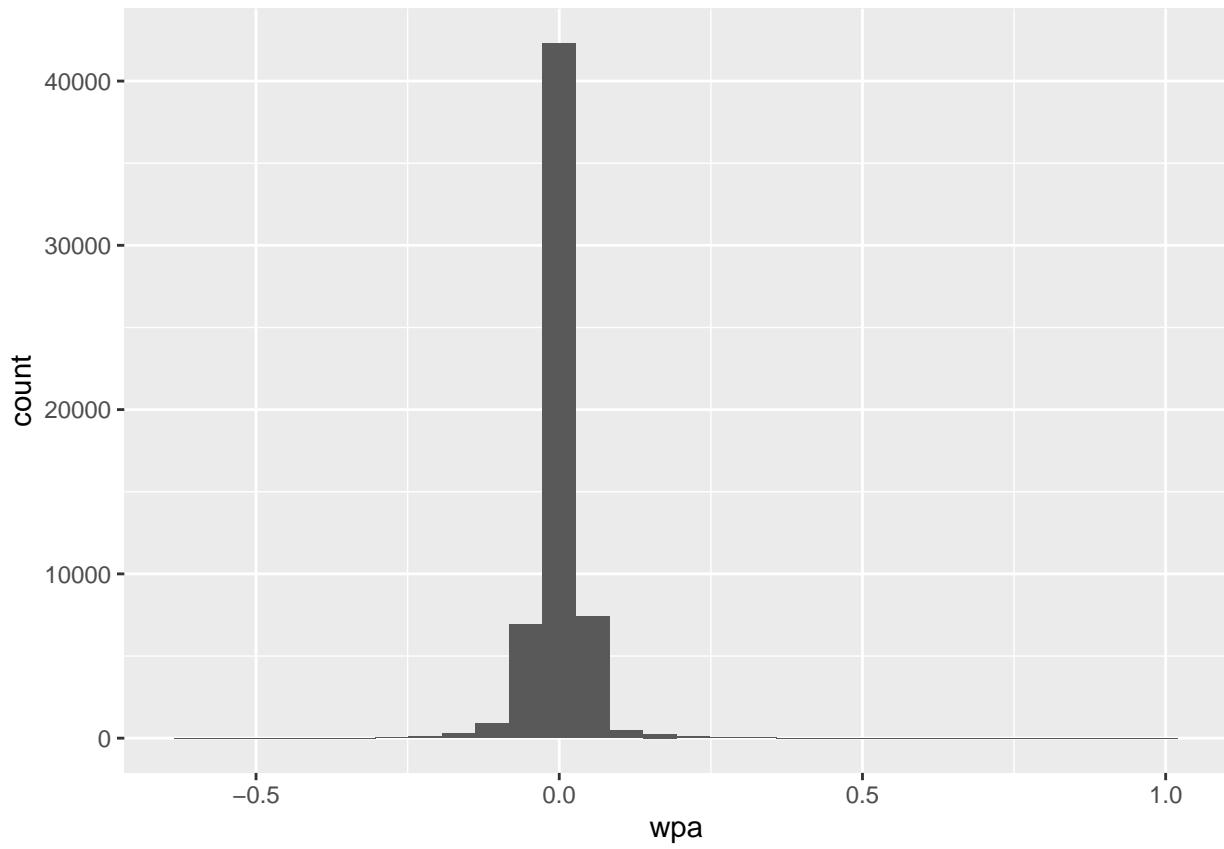
```
ggplot(data_main, aes(x=year)) + geom_line(aes(fill=..count.., color = play_type), stat="bin", binwidth=1) +
  scale_x_continuous(breaks = c(1999:2021), labels = c(1999:2021), limits = c(1999, 2021)) +
  theme_bw() +
  theme(axis.text.x = element_text(angle = 45, hjust=1),
        panel.grid.minor.x = element_blank()) +
  labs(title = "Count of NFL 4th Down Plays",
       subtitle = "4th Down Plays Between 1999 and 2022",
       x = "Year", y = "# of Plays",
       color = "Decision")
```

## Count of NFL 4th Down Plays

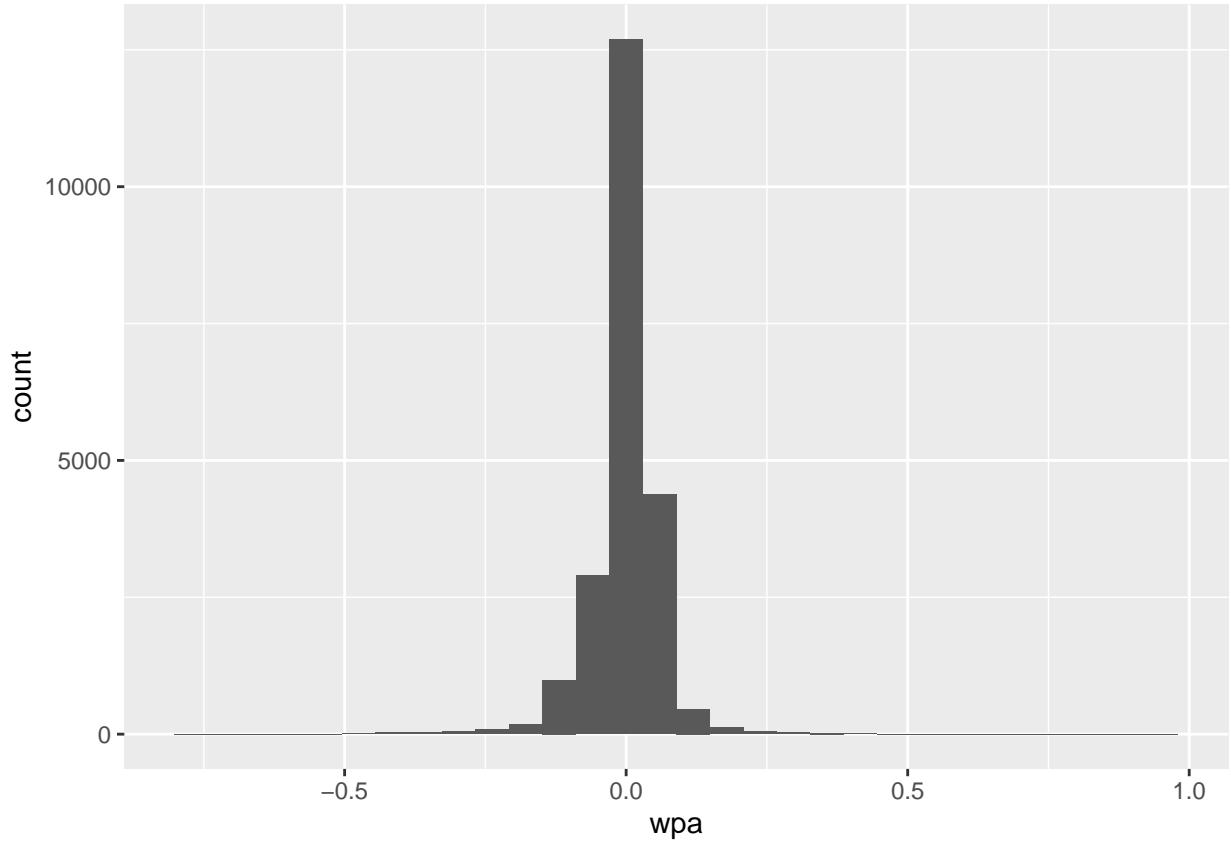
### 4th Down Plays Between 1999 and 2022



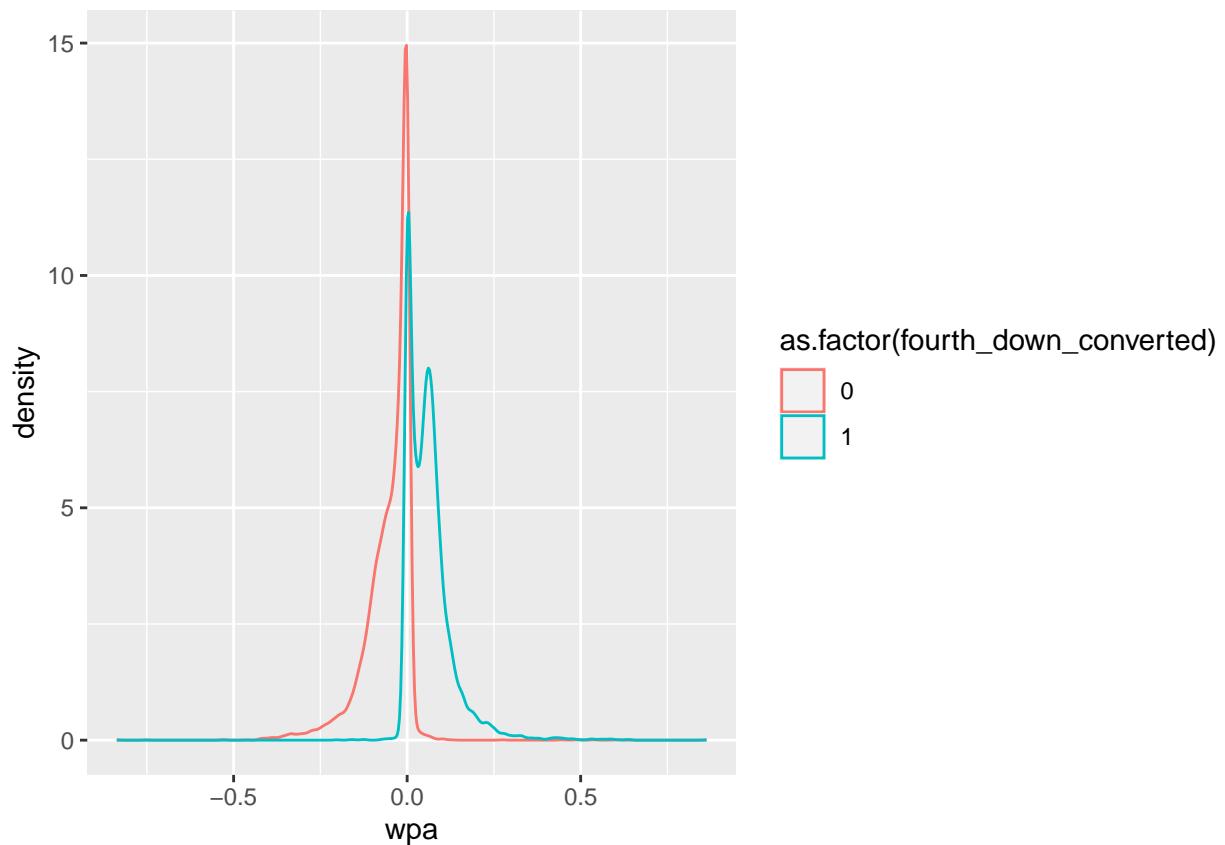
```
ggplot(punt_data, aes(x = wpa)) +  
  geom_histogram()
```



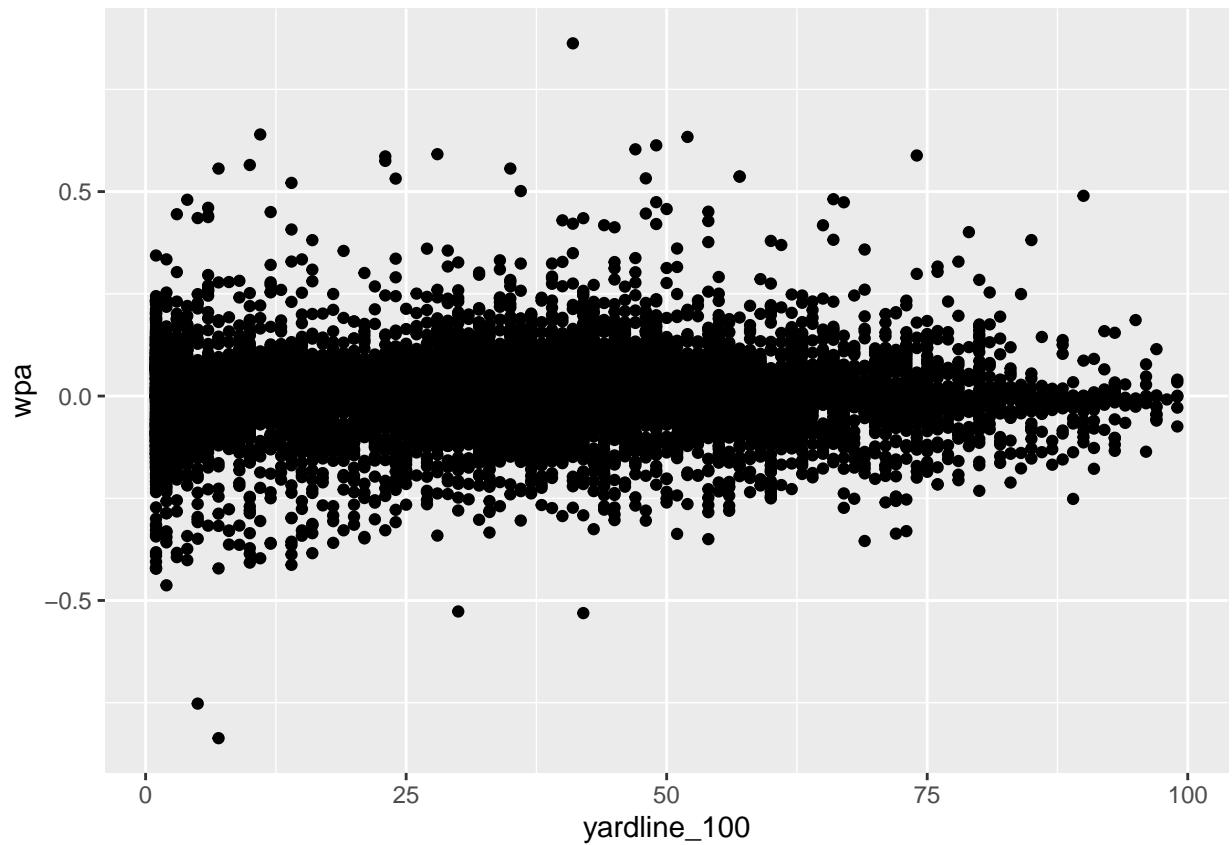
```
ggplot(kick_data, aes(x = wpa)) +  
  geom_histogram()
```



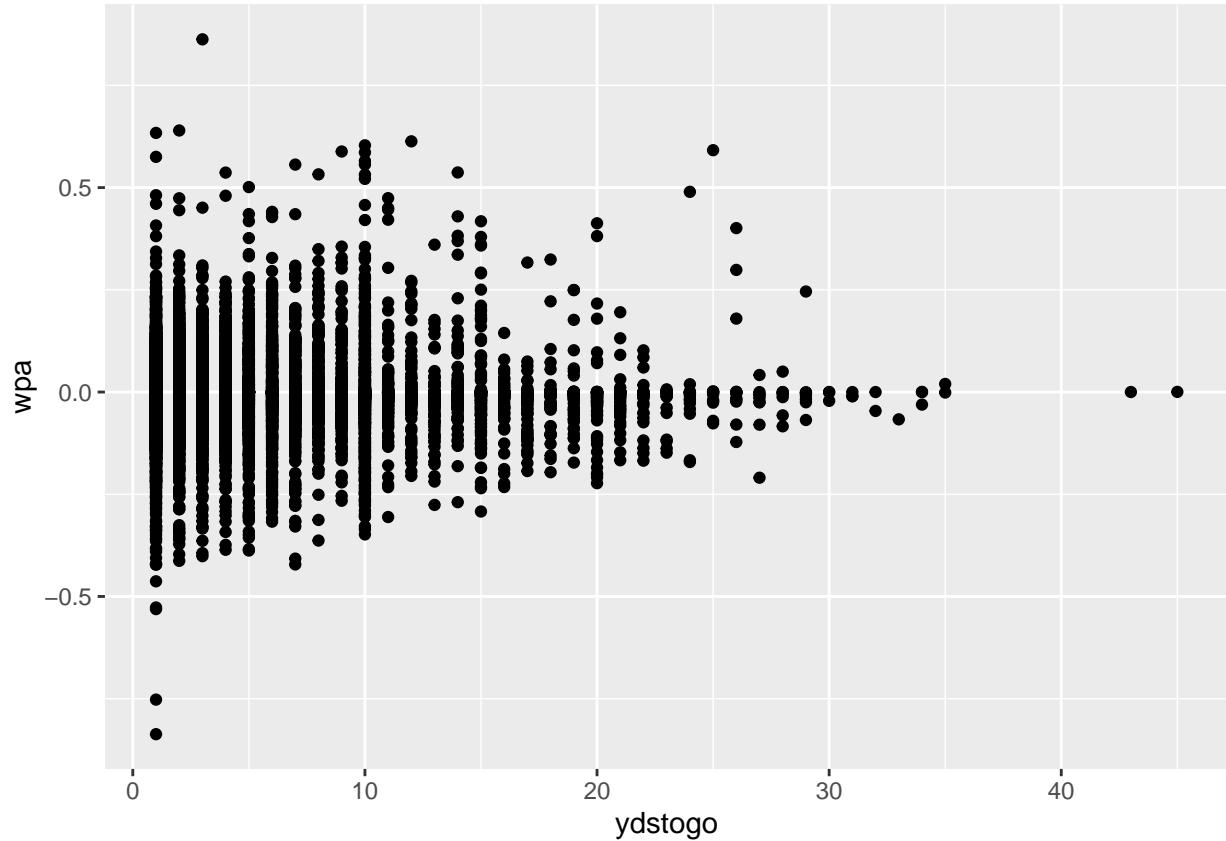
```
ggplot(go_data, aes(x = wpa, color = as.factor(fourth_down_converted))) +  
  geom_density()
```



```
ggplot(go_data, aes(x = yardline_100, y = wpa)) +  
  geom_point()
```



```
ggplot(go_data, aes(x = ydstogo, y = wpa)) +  
  geom_point()
```



```
data_main %>%
  group_by(play_type) %>%
  summarise(mean_wpa = mean(wpa)) %>%
  kable(digits = 5)
```

play_type	mean_wpa
field_goal	-0.00037
go	0.00396
punt	-0.00044

## Modeling

The goal with our models is to create the best possible models for predicting WPA (win probability added) given a certain game state.

### Punt Model

```
set.seed(12321)
sample <- sample(c(TRUE, FALSE), nrow(punt_data), replace = TRUE, prob = c(0.7, 0.3))
punt_data_train <- punt_data[sample, ]
```

```

punt_data_test <- punt_data[!sample, ]
punt_train_x <- data.matrix(punt_data_train[, -1])
punt_train_y <- punt_data_train$wpa
punt_test_x <- data.matrix(punt_data_test[, -1])
punt_test_y <- punt_data_test$wpa
punt_xgb_train = xgb.DMatrix(data = punt_train_x, label = punt_train_y)
punt_xgb_test = xgb.DMatrix(data = punt_test_x, label = punt_test_y)
punt_watchlist = list(train=punt_xgb_train, test=punt_xgb_test)
punt_model = xgb.train(data = punt_xgb_train, max_depth = 8, colsample_bylevel = 1, colsample_bytree =
punt_model_xgboost = xgboost(data = punt_xgb_train, max_depth = 8, colsample_bylevel = 1, colsample_by
punt_pred_y = predict(punt_model_xgboost, punt_xgb_test)
print(paste("XG Boost RMSE:", RMSE(punt_test_y, punt_pred_y)))

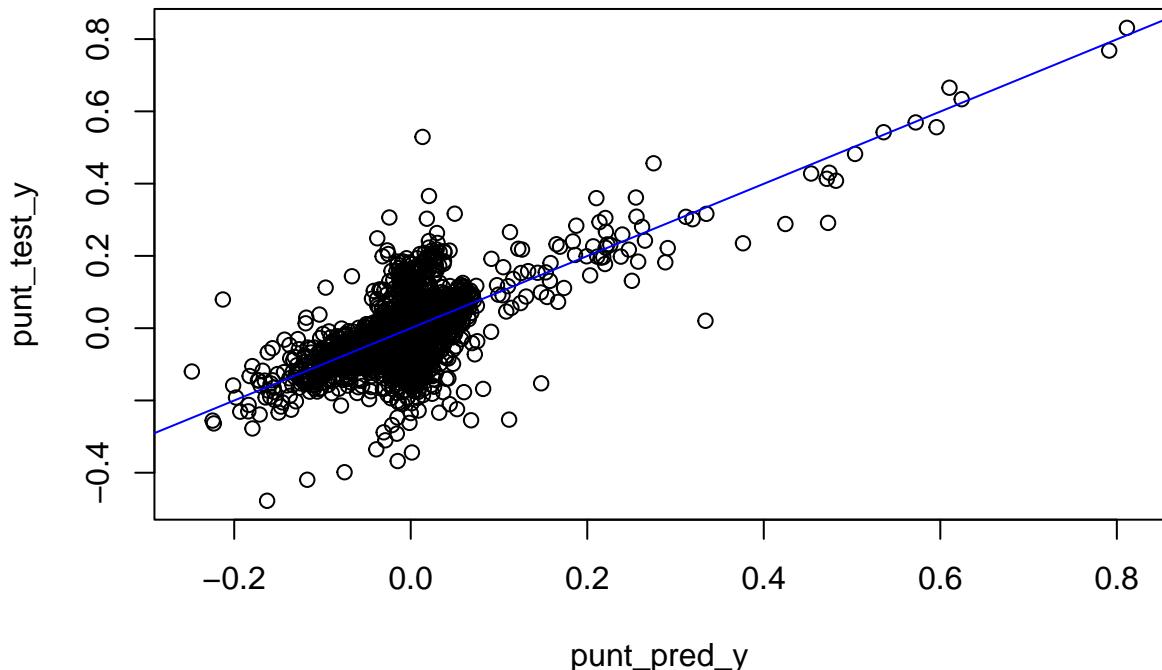
```

```
## [1] "XG Boost RMSE: 0.0294369905331241"
```

```

plot(y=punt_test_y, x=punt_pred_y)
abline(lm(punt_test_y ~ punt_pred_y), col = "blue")

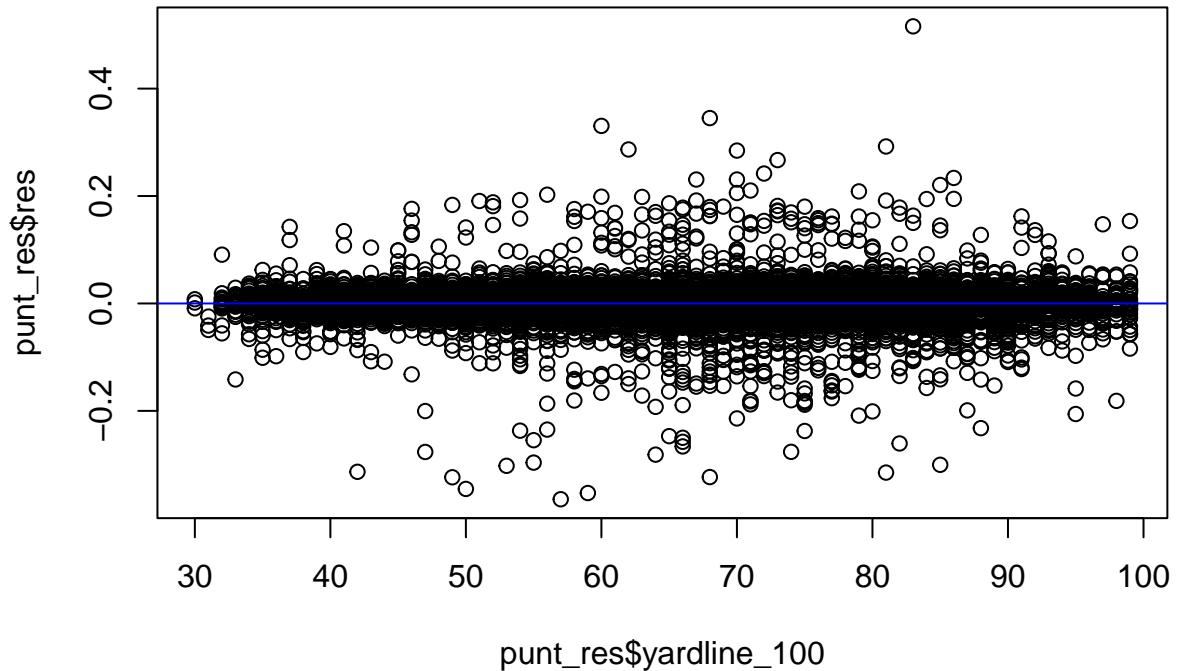
```



```

punt_res = data.frame(test_y= punt_test_y, pred_y = punt_pred_y, res = (punt_test_y-punt_pred_y), punt_
plot(punt_res$yardline_100, punt_res$res)
abline(0,0, col="blue")

```



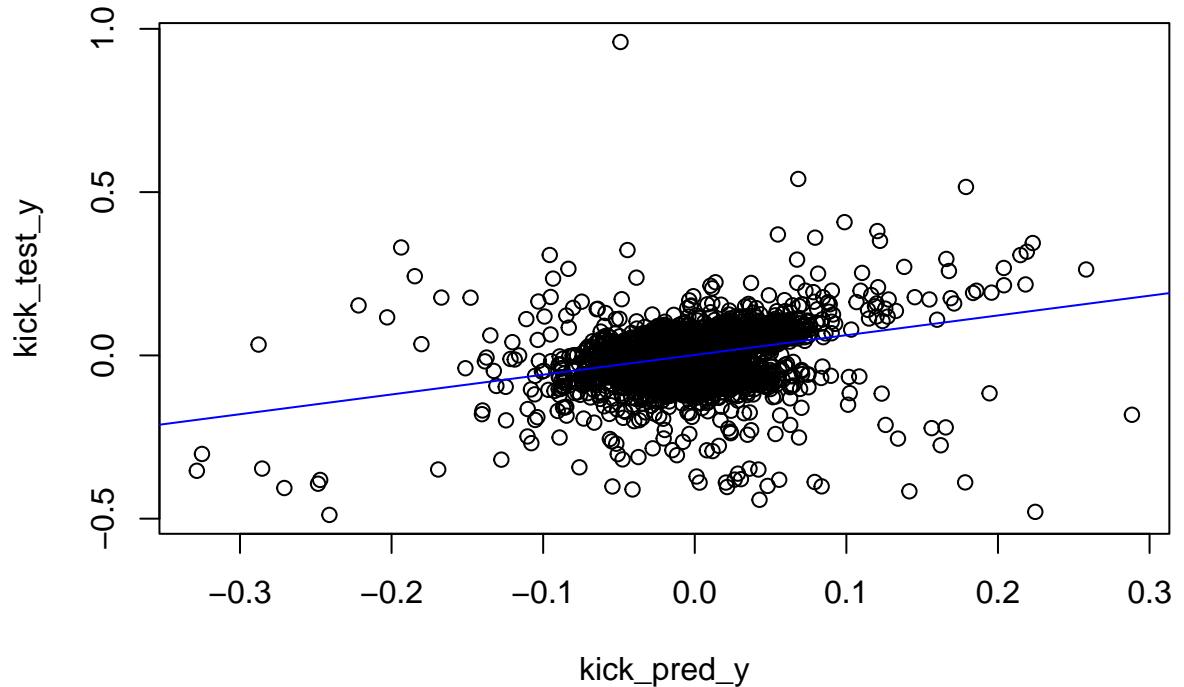
Our optimal punt XG Boost model has a RMSE of 0.0216, which is the lowest RMSE we obtained throughout our XG Boost, Lasso, and Ridge Regression modelling process. Therefore, we will continue with this as our punt WPA model.

## Kick Model

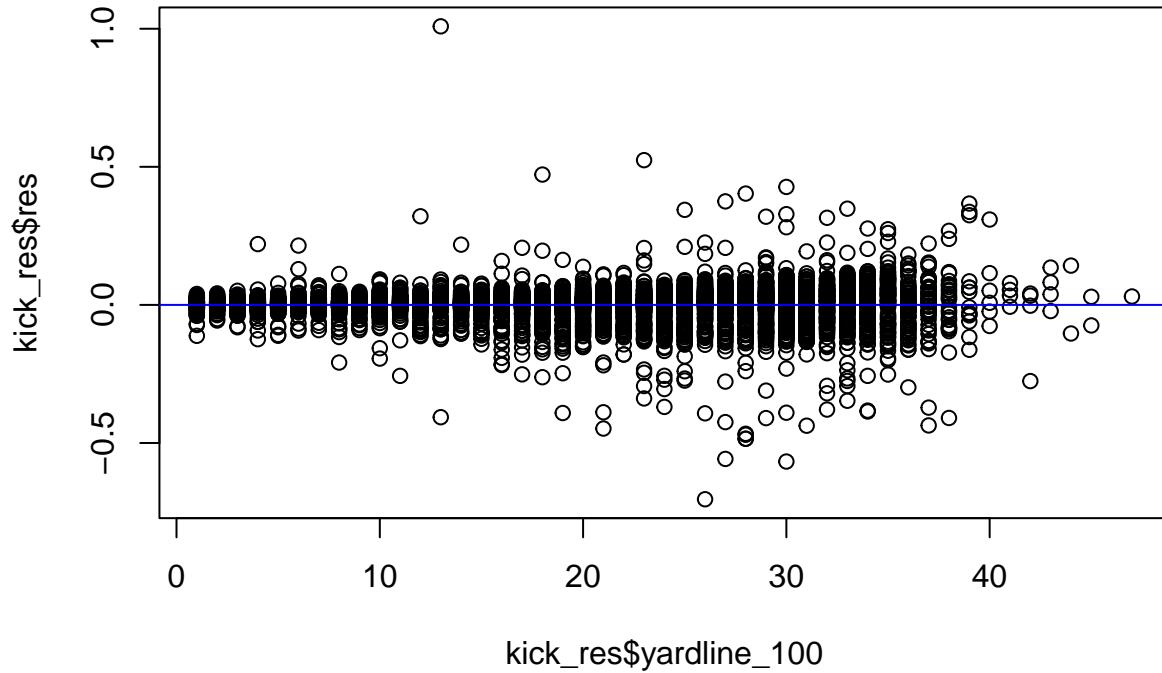
```
set.seed(12321)
sample <- sample(c(TRUE, FALSE), nrow(kick_data), replace = TRUE, prob = c(0.7, 0.3))
kick_data_train <- kick_data[sample, ]
kick_data_test <- kick_data[!sample, ]
kick_train_x <- data.matrix(kick_data_train[, -1])
kick_train_y <- kick_data_train$wpa
kick_test_x <- data.matrix(kick_data_test[, -1])
kick_test_y <- kick_data_test$wpa
kick_xgb_train = xgb.DMatrix(data = kick_train_x, label = kick_train_y)
kick_xgb_test = xgb.DMatrix(data = kick_test_x, label = kick_test_y)
kick_watchlist = list(train=kick_xgb_train, test=kick_xgb_test)
kick_model = xgb.train(data = kick_xgb_train, max_depth = 6, colsample_bylevel = 1, colsample_bytree =
kick_model_xgboost = xgboost(data = kick_xgb_train, max_depth = 6, colsample_bylevel = 1, colsample_by
kick_pred_y = predict(kick_model_xgboost, kick_xgb_test)
print(paste("XG Boost RMSE:", RMSE(kick_test_y, kick_pred_y)))

## [1] "XG Boost RMSE: 0.0598292868845007"
```

```
plot(y=kick_test_y, x=kick_pred_y)
abline(lm(kick_test_y ~ kick_pred_y), col = "blue")
```



```
kick_res = data.frame(test_y= kick_test_y, pred_y = kick_pred_y, res = (kick_test_y-kick_pred_y), kick_100 = yardline_100)
plot(kick_res$yardline_100, kick_res$res)
abline(0,0, col="blue")
```



Our optimal kick XG Boost model has a RMSE of 0.3921, which is the lowest RMSE we obtained throughout our XG Boost, Lasso, and Ridge Regression modelling process. Therefore, we will continue with this as our kick WPA model.

## Go Model

```

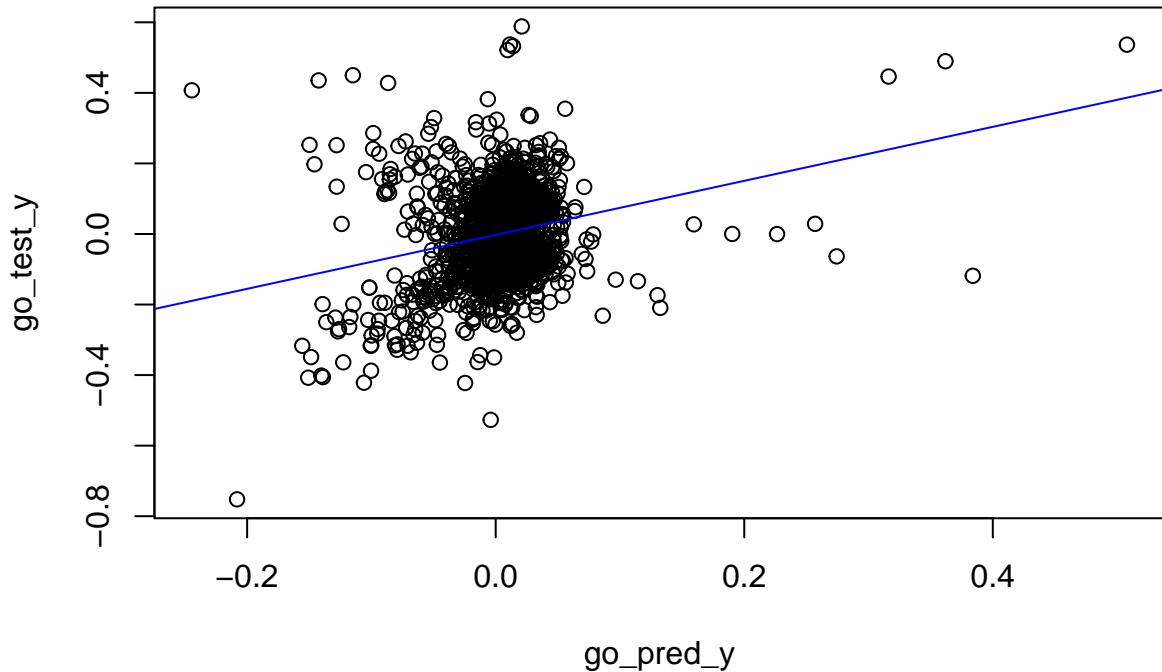
go_data <- go_data %>%
  select(-fourth_down_converted)

set.seed(12321)
sample <- sample(c(TRUE, FALSE), nrow(go_data), replace = TRUE, prob = c(0.7, 0.3))
go_data_train <- go_data[sample, ]
go_data_test <- go_data[!sample, ]
go_train_x <- data.matrix(go_data_train[, -1])
go_train_y <- go_data_train$wpa
go_test_x <- data.matrix(go_data_test[, -1])
go_test_y <- go_data_test$wpa
go_xgb_train = xgb.DMatrix(data = go_train_x, label = go_train_y)
go_xgb_test = xgb.DMatrix(data = go_test_x, label = go_test_y)
go_watchlist = list(train=go_xgb_train, test=go_xgb_test)
go_model = xgb.train(data = go_xgb_train, max_depth = 8, colsample_bytree = 1, colsample_bylevel = 1, nrounds = 100)
go_model_xgboost = xgboost(data = go_xgb_train, max_depth = 6, colsample_bytree = 1, colsample_bylevel = 1, nrounds = 100)
go_pred_y = predict(go_model_xgboost, go_xgb_test)
print(paste("XG Boost RMSE:", RMSE(go_test_y, go_pred_y)))

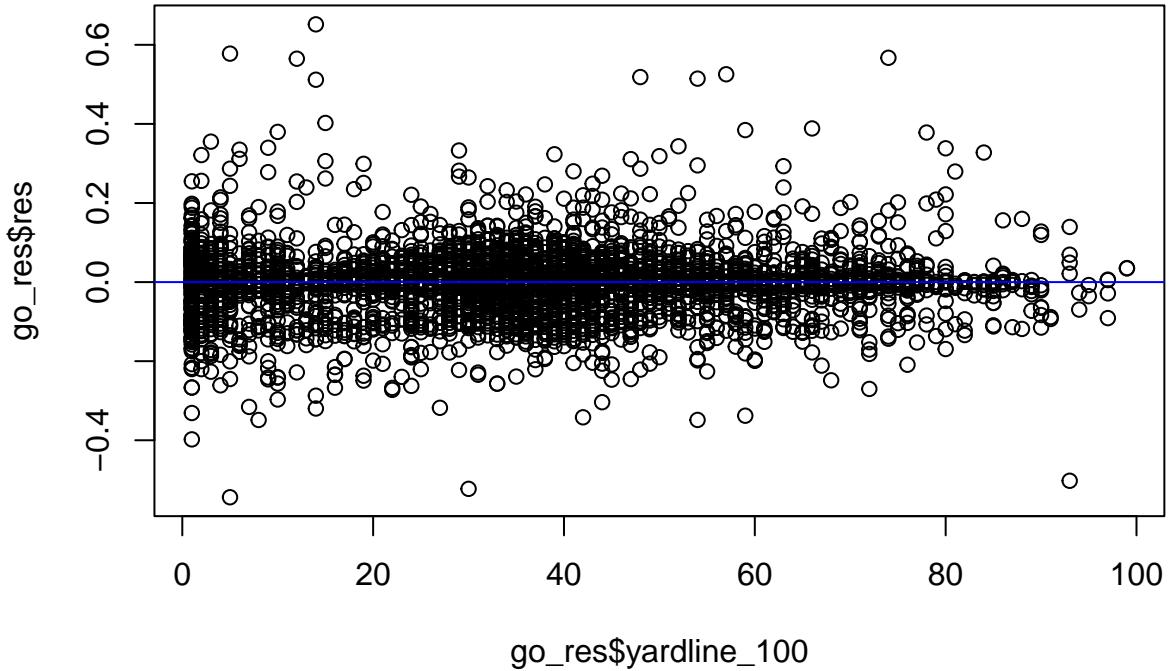
```

```
## [1] "XG Boost RMSE: 0.0867906629769721"
```

```
plot(y=go_test_y, x=go_pred_y)
abline(lm(go_test_y ~ go_pred_y), col = "blue")
```



```
go_res = data.frame(test_y= go_test_y, pred_y = go_pred_y, res = (go_test_y-go_pred_y), go_test_x)
plot(go_res$yardline_100, go_res$res)
abline(0,0, col="blue")
```



Our optimal go for it XG Boost model has a RMSE of 0.5492, which is the lowest RMSE we obtained throughout our XG Boost, Lasso, and Ridge Regression modelling process. Therefore, we will continue with this as our go for it WPA model.

## Results

```

pred_data <- data_main %>%
  select(-c(down, fourth_down_converted, fourth_down_failed, ep, epa, year,
           week, posteam, play_type, posteam_type)) %>%
  select(wpa, everything())

data_x <- data.matrix(pred_data[, -1])
data_y <- pred_data$wpa
xgb_matrix = xgb.DMatrix(data = data_x, label = data_y)

ewpa_punt = predict(punt_model_xgboost, xgb_matrix)
ewpa_kick = predict(kick_model_xgboost, xgb_matrix)
ewpa_go = predict(go_model_xgboost, xgb_matrix)

pred_data <- pred_data %>%
  mutate(yard_line = ifelse(yardline_100 > 50, -(100-yardline_100), yardline_100),
         game_minutes_remaining = paste0(trunc(game_seconds_remaining/60), ":", game_seconds_remaining%))

pred_data = cbind(pred_data, play_type = data_main$play_type, team = data_main$posteam, year = data_main$year)
  
```

```

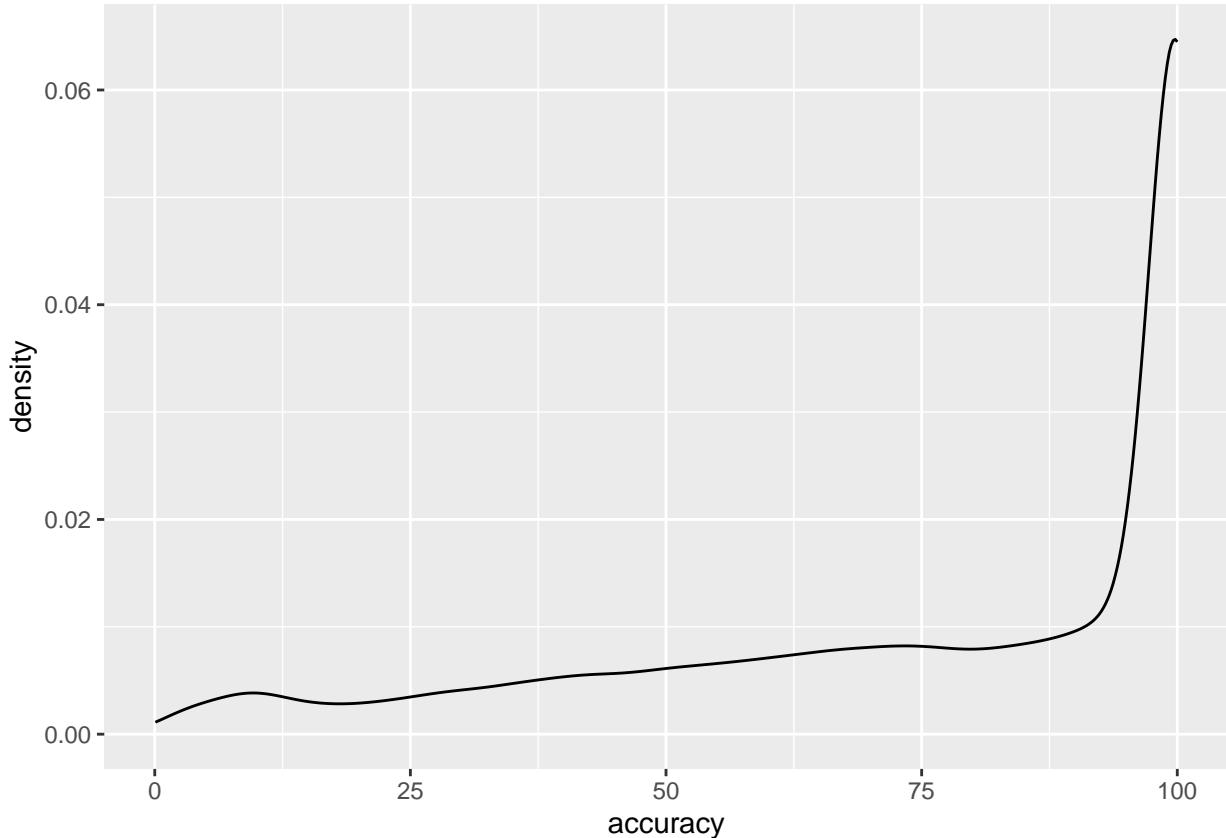
mutate(
  #ewpa_punt = ifelse(yard_line < 31, -1, ewpa_punt),
  recommendation = case_when((ewpa_punt > ewpa_kick) & (ewpa_punt > ewpa_go) ~ "punt",
                             (ewpa_kick > ewpa_punt) & (ewpa_kick > ewpa_go) ~ "field_goal",
                             (ewpa_go > ewpa_kick) & (ewpa_go > ewpa_punt) ~ "go"),
  correct_call = ifelse(play_type == recommendation, "Yes", "No"),
  ewpa_max = pmax(ewpa_punt, ewpa_kick, ewpa_go),
  ewpa_play = case_when((play_type == "punt") ~ ewpa_punt,
                        (play_type == "field_goal") ~ ewpa_kick,
                        (play_type == "go") ~ ewpa_go),
  execution = wpa - ewpa_max)

pred_data$ewpa_max_percentile <- ecdf(pred_data$ewpa_max)(pred_data$ewpa_max)
pred_data$ewpa_play_percentile <- ecdf(pred_data$ewpa_max)(pred_data$ewpa_play)
pred_data$accuracy <- 100 * (1 - (pred_data$ewpa_max_percentile - pred_data$ewpa_play_percentile))

#pred_data %>%
#  filter(correct_call == "No") %>%
#  select(yard_line, game_minutes_remaining, ydstogo, score_differential, play_type, recommendation)

#example <- confusionMatrix(data=predicted_value, reference = expected_value)
#example
ggplot(pred_data, aes(x = accuracy)) +
  geom_density()

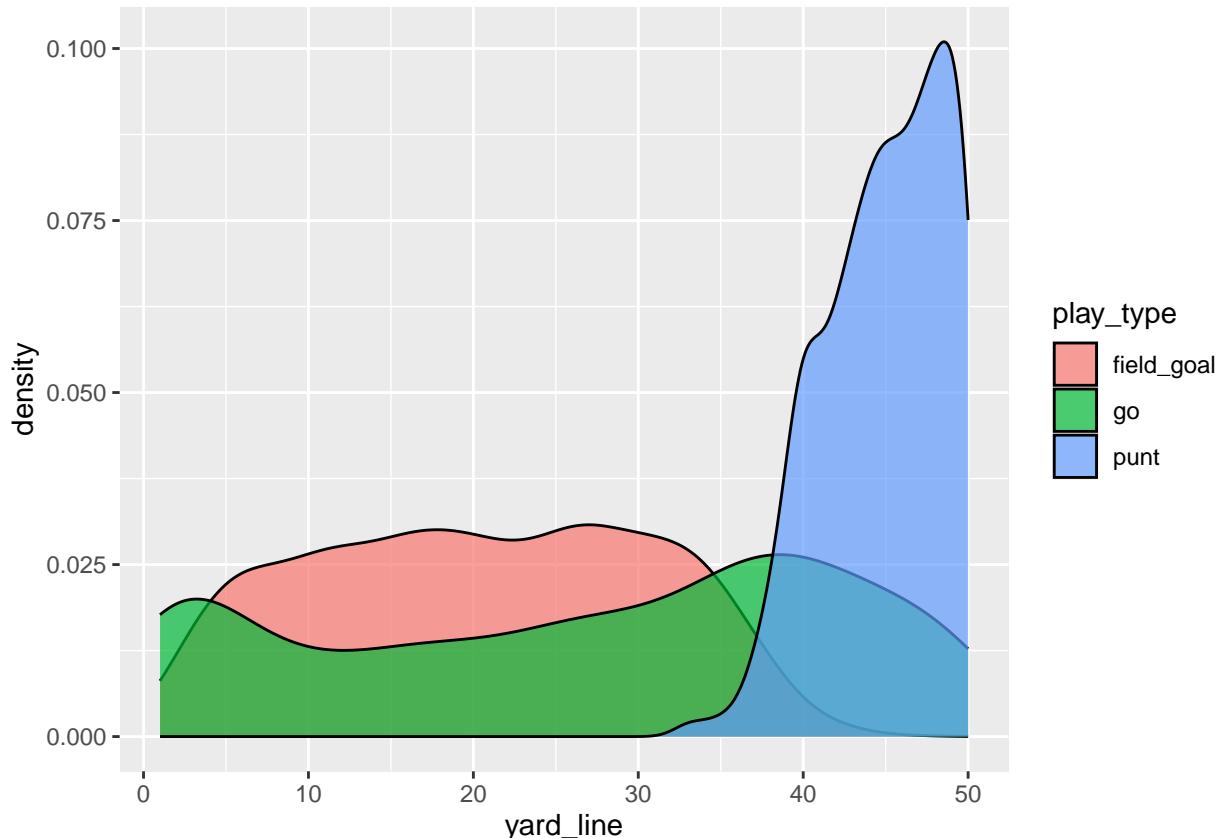
```



```

pred_data %>%
  filter(yard_line>0,
        year > 2020) %>%
  ggplot(aes(x=yard_line, fill = play_type)) + geom_density(alpha=0.7)

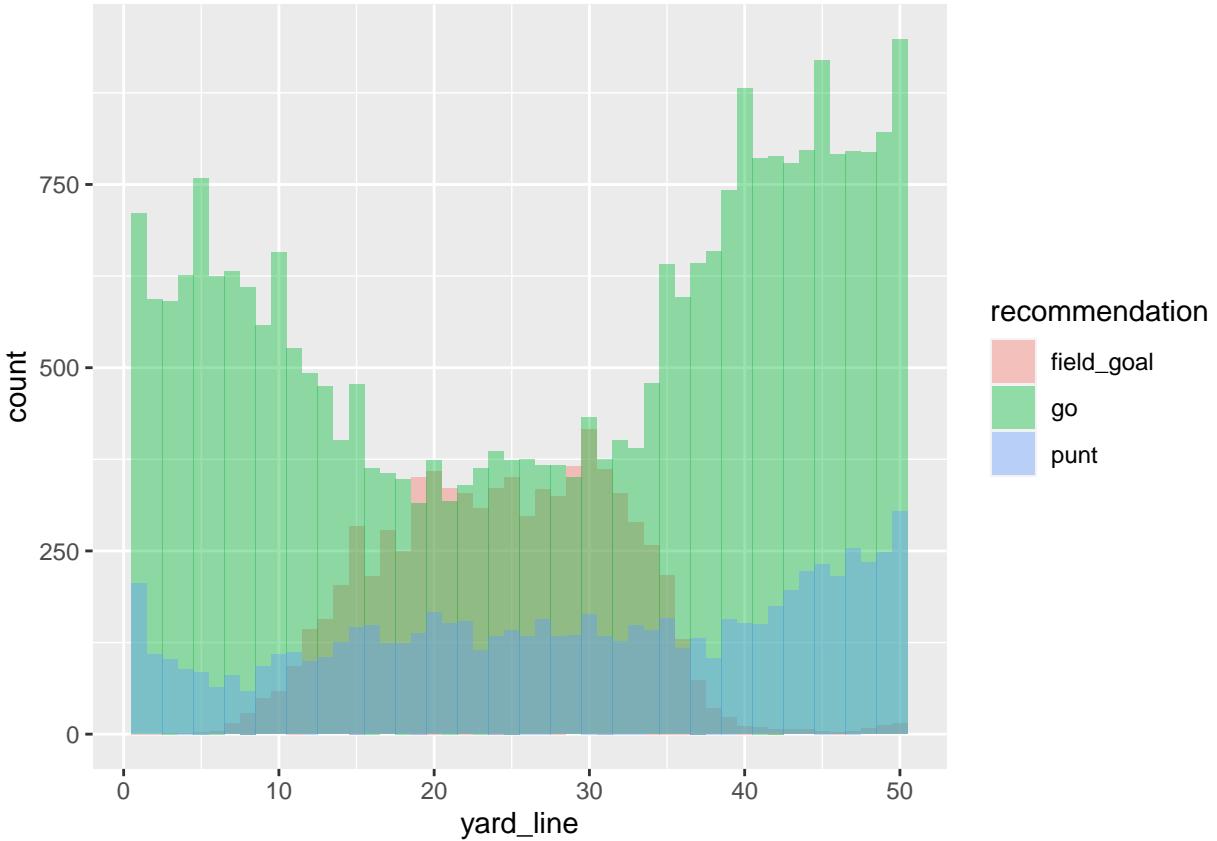
```



```

pred_data %>%
  filter(yard_line>0) %>%
  ggplot(aes(x=yard_line, fill = recommendation)) + geom_histogram(binwidth = 1, alpha=.4, position="identity")

```



```

crucial_data <- data_main %>%
  select(-c(down, fourth_downConverted, fourth_downFailed, ep, epa, year,
    week, posteam, play_type, posteam_type)) %>%
  select(wpa, everything()) %>%
  filter((yardline_100 < 51 | ydstogo < 5) & (abs(score_differential) < 14.5) & (wp > 0.1))

data_x <- data.matrix(crucial_data[, -1])
data_y <- crucial_data$wpa
xgb_matrix = xgb.DMatrix(data = data_x, label = data_y)

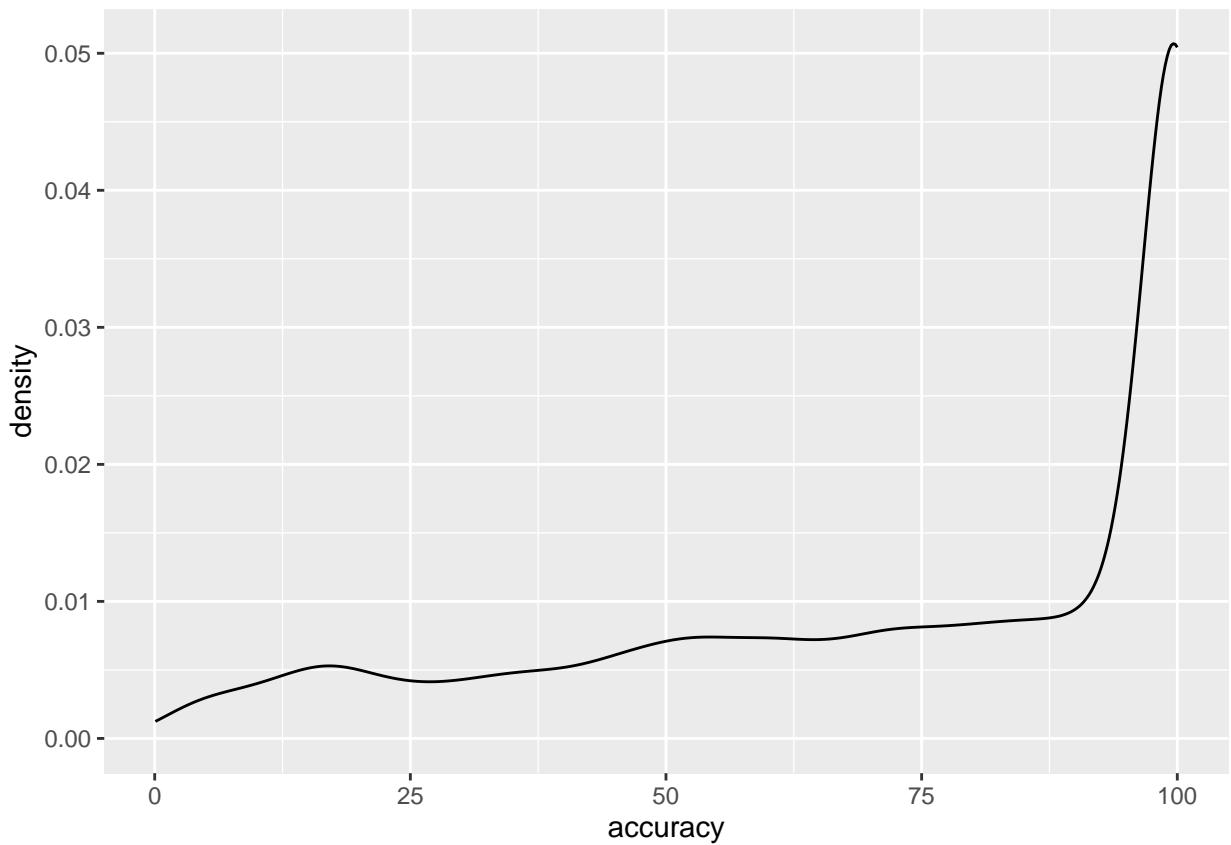
ewpa_punt = predict(punt_model_xgboost, xgb_matrix)
ewpa_kick = predict(kick_model_xgboost, xgb_matrix)
ewpa_go = predict(go_model_xgboost, xgb_matrix)

filtered_data <- data_main %>%
  filter((yardline_100 < 51 | ydstogo < 5) & (abs(score_differential) < 14.5) & (wp > 0.1))

crucial_data <- cbind(crucial_data, play_type = filtered_data$play_type,
  team = filtered_data$posteam, year = filtered_data$year, ewpa_punt, ewpa_kick, ewpa_go) %>%
  mutate(recommendation = case_when((ewpa_punt > ewpa_kick) & (ewpa_punt > ewpa_go) ~ "punt",
    (ewpa_kick > ewpa_punt) & (ewpa_kick > ewpa_go) ~ "field_goal",
    (ewpa_go > ewpa_kick) & (ewpa_go > ewpa_punt) ~ "go"),
  correct_call = ifelse(play_type == recommendation, "Yes", "No"),
  ewpa_max = pmax(ewpa_punt, ewpa_kick, ewpa_go),
  ewpa_play = case_when((play_type == "punt") ~ ewpa_punt,
    (play_type == "field_goal") ~ ewpa_kick,
    (play_type == "go") ~ ewpa_go))

```

```
(play_type == "go") ~ ewpa_go),  
execution = wpa - ewpa_max)  
  
crucial_data$ewpa_max_percentile <- ecdf(crucial_data$ewpa_max)(crucial_data$ewpa_max)  
crucial_data$ewpa_play_percentile <- ecdf(crucial_data$ewpa_max)(crucial_data$ewpa_play)  
crucial_data$accuracy <- 100 * (1 - (crucial_data$ewpa_max_percentile - crucial_data$ewpa_play_percentile))  
  
ggplot(crucial_data, aes(x = accuracy)) +  
  geom_density()
```



```
crucial_data %>%  
  arrange(desc(year))
```

```

##          wpa yardline_100 game_seconds_remaining qtr ydstogo
## 1: -0.002238333      51                  3296   1      1
## 2: -0.019976735       6                  2899   1      3
## 3: -0.018679649      66                  2820   1      1
## 4: -0.095877826      27                  2455   2      4
## 5:  0.030603722      27                  1828   2      5
##    ---
## 46803:  0.028520644     11                  2369   2     10
## 46804: -0.044980898     10                  1819   2     10
## 46805: -0.073177233     29                  1800   3      5
## 46806:  0.060336888     57                  876    4      1

```

```

## 46807: 0.002091736          25          377    4     7
##           score_differential no_score_prob opp_fg_prob opp_safety_prob
## 1:                      0 0.011813287 0.177440062 1.565513e-03
## 2:                      0 0.000751453 0.003594615 5.229356e-05
## 3:                     -3 0.046969537 0.196949929 3.269824e-03
## 4:                     -3 0.027033767 0.055807115 7.228962e-04
## 5:                    -10 0.198925869 0.031771343 7.081460e-05
## ---
## 46803:                   3 0.012635923 0.010552165 1.447756e-04
## 46804:                   6 0.058551728 0.001686924 1.439996e-05
## 46805:                  -9 0.002216538 0.084979071 1.469645e-03
## 46806:                 -10 0.084862515 0.183874503 2.234099e-03
## 46807:                  -3 0.078442221 0.041841515 6.270266e-04
##           opp_td_prob   fg_prob   safety_prob      td_prob      wp
## 1: 0.3223033249 0.1836291 9.829091e-03 0.2934196293 0.4737365
## 2: 0.0052889088 0.9839867 4.810112e-05 0.0062778788 0.5815628
## 3: 0.2704848647 0.1542629 3.105036e-03 0.3249579072 0.4060214
## 4: 0.0902073976 0.7886681 3.824418e-04 0.0371782679 0.4443144
## 5: 0.0026942235 0.7643739 1.263318e-04 0.0020374833 0.2256564
## ---
## 46803: 0.0191436318 0.9411797 3.276049e-04 0.0160162515 0.5929926
## 46804: 0.0005263558 0.9385480 1.056186e-04 0.0005670034 0.7959621
## 46805: 0.1227084289 0.7049640 9.205743e-04 0.0827417670 0.2314794
## 46806: 0.3207639158 0.1546367 6.023312e-03 0.2476048768 0.1292370
## 46807: 0.0705068382 0.7703825 5.302751e-04 0.0376695779 0.4555477
##           posteam_timeouts_remaining defteam_timeouts_remaining play_type team
## 1:                      2                      3       punt  BAL
## 2:                      2                      3 field_goal  BAL
## 3:                      3                      2       punt  NYJ
## 4:                      3                      2 field_goal  NYJ
## 5:                      2                      0 field_goal  NYJ
## ---
## 46803:                   3                      3 field_goal  LA
## 46804:                   1                      3 field_goal  LA
## 46805:                   3                      3 field_goal  TEN
## 46806:                   2                      3       go  TEN
## 46807:                   1                      3 field_goal  TEN
##           year      ewpa_punt      ewpa_kick      ewpa_go recommendation correct_call
## 1: 2022 -0.006766066 -0.0462302752 0.036177561       go        No
## 2: 2022 -0.003692099 -0.0249907244 0.132701740       go        No
## 3: 2022 -0.016130652 -0.0469167158 0.022320151       go        No
## 4: 2022 -0.042633358 -0.0612212941 0.020426100       go        No
## 5: 2022  0.009661306  0.0225406438 0.158307135       go        No
## ---
## 46803: 1999 -0.014442386  0.0231197570 0.108514041       go        No
## 46804: 1999 -0.010340491 -0.0057053478 0.099041216       go        No
## 46805: 1999 -0.012865994 -0.0595337227 0.007767885       go        No
## 46806: 1999 -0.003388926 -0.0826969445 0.026677499       go       Yes
## 46807: 1999 -0.079701655 -0.0002697437 -0.020439010 field_goal       Yes
##           ewpa_max      ewpa_play      execution ewpa_max_percentile
## 1: 0.0361775607 -0.0067660660 -0.038415894 0.7226911
## 2: 0.1327017397 -0.0249907244 -0.152678475 0.8926870
## 3: 0.0223201513 -0.0161306523 -0.040999800 0.5309035
## 4: 0.0204261001 -0.0612212941 -0.116303926 0.4971906

```

```

##      5:  0.1583071351  0.0225406438 -0.127703413           0.9269767
##    ---
## 46803:  0.1085140407  0.0231197570 -0.079993397           0.8544448
## 46804:  0.0990412161 -0.0057053478 -0.144022115           0.8418185
## 46805:  0.0077678850 -0.0595337227 -0.080945118           0.2660286
## 46806:  0.0266774986  0.0266774986  0.033659389           0.6137330
## 46807: -0.0002697437 -0.0002697437  0.002361479           0.1477557
##      ewpa_play_percentile accuracy
##      1:          0.078428440 35.57374
##      2:          0.013929540 12.12425
##      3:          0.030914180 50.00107
##      4:          0.001837332 50.46467
##      5:          0.534813169 60.78364
##    ---
## 46803:          0.545409875 69.09650
## 46804:          0.087743286 24.59248
## 46805:          0.001986882 73.59583
## 46806:          0.613732989 100.00000
## 46807:          0.147755678 100.00000

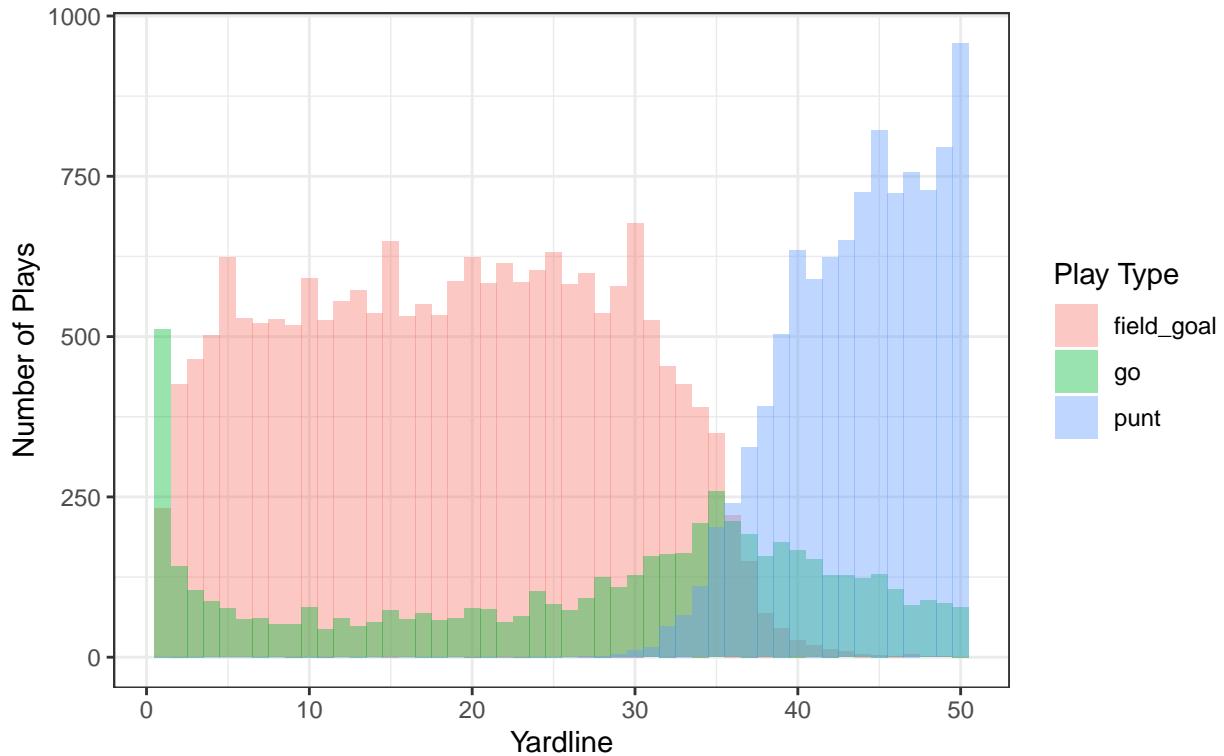
top_decisions <- crucial_data %>%
  arrange(desc(execution)) %>%
  select(ewpa_punt, ewpa_kick, ewpa_go, recommendation, play_type, wpa, execution, accuracy) %>%
  head(10)

crucial_data %>%
  filter(yardline_100<51) %>%
  ggplot(aes(x=yardline_100, fill = play_type)) + geom_histogram(binwidth = 1, alpha=.4, position="identity")
  theme_bw()+
  labs(title = "Real Life Decisions", subtitle = "1999-2022, opponents' side of field, game within 14 p"

```

## Real Life Decisions

1999–2022, opponents' side of field, game within 14 points, win probability > 0.1

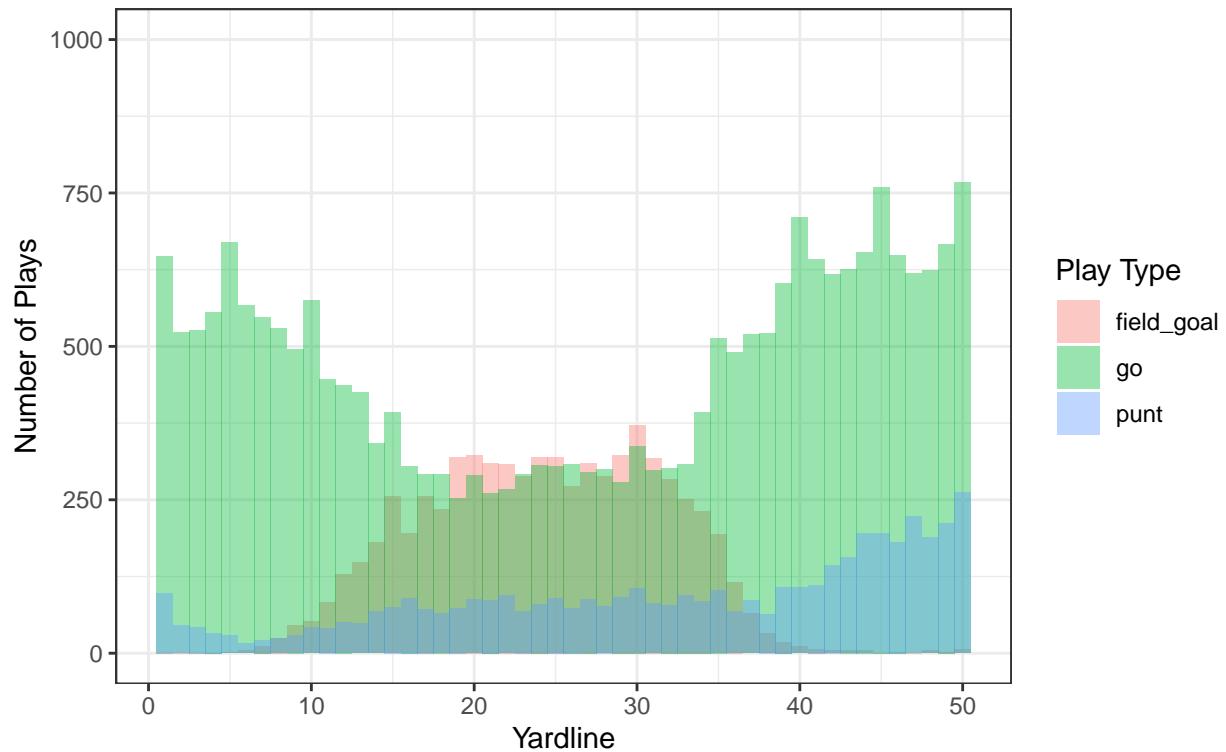


```
crucial_data %>%
  filter(yardline_100<51) %>%
  ggplot(aes(x=yardline_100, fill = recommendation)) + geom_histogram(binwidth = 1, alpha=.4, position=
  theme_bw()+
  scale_y_continuous(limits = c(0,1000))+
```

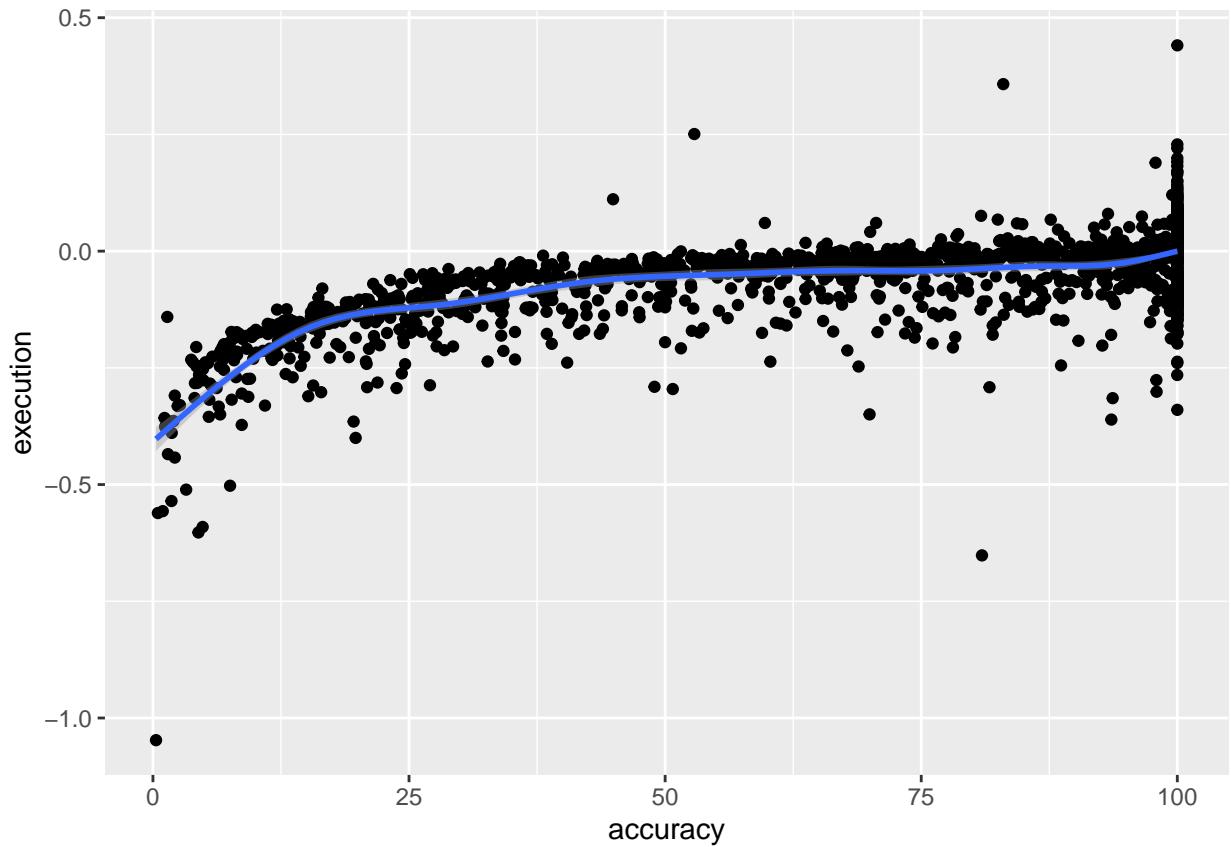
labs(title = "Model Recommended Decisions", subtitle = "1999–2022, opponents' side of field, game within 14 points, win probability > 0.1")

## Model Recommended Decisions

1999–2022, opponents' side of field, game within 14 points, win probability > 0.1



```
ggplot(crucial_data[year == 2021], aes(x = accuracy, y = execution)) +  
  geom_point() +  
  geom_smooth()
```



```
top_decisions %>%
  kable(digits = 3)
```

ewpa_punt	ewpa_kick	ewpa_go	recommendation	play_type	wpa	execution	accuracy
-0.142	-0.091	0.014	go	go	0.532	0.518	100.000
0.013	-0.080	-0.003	punt	punt	0.529	0.516	100.000
-0.111	-0.258	0.002	go	go	0.481	0.479	100.000
-0.024	-0.194	-0.115	punt	go	0.450	0.474	98.494
-0.024	-0.312	0.115	go	go	0.586	0.471	100.000
0.071	-0.093	0.010	punt	go	0.521	0.450	48.681
-0.082	-0.247	0.033	go	go	0.474	0.441	100.000
-0.005	-0.336	-0.142	punt	go	0.435	0.441	90.826
-0.126	-0.145	0.061	go	go	0.474	0.413	100.000
-0.069	-0.361	0.012	go	go	0.418	0.406	100.000

## Discussion and Conclusion